

STP598-Assignment 4

Hao Wang
hwang306

November 14, 2017

1 Question 1

Create an R function named `RLMstep` that performs (hybrid) stepwise model selection based on the method illustrated in p. 11 of the article entitled 'Robust Stepwise Regression'. The function should be based on the existing `rlm` function; use the default settings. The function should take as input: 1. a column vector y (response) 2. a matrix X of predictors (predictors should be in columns)

Your function should report: a) the predictors that are included in the final robust regression model, and b) coefficient estimates, and their BCa confidence intervals

The hybrid method of stepwise regression considers both forward and backwards directions. Similar to forward stepwise regression, in a hybrid model, variables are added sequentially. However, after the new variable is added, this method also removes variables that are no longer statistically significant in the regression. The hybrid method mimics the power of best subset regression, but also has the computational advantages of forward and backwards stepwise regression.

My function is based on the method from the Agostinelli (2002)'s idea of robust stepwise regression. I use the function `wle.stepwise` in the `wle` package to perform the variable selection based on weighted stepwise regression (which is the package supplement of the article).

The function `RLMstep` will report the predictors that are included, as well as the coefficients from the `rlm` function. I use the Boston data from `MASS` package to display as an example.

In the example, out of the 14 variables, 3 variables are selected: `crim`, `zn` and `indus`. Coefficients are reported.

```
library(MASS)
library(wle) #This package reflects the paper Agostinelli 2002
```

```
## Loading required package: circular
##
## Attaching package: 'circular'
## The following objects are masked from 'package:stats':
##
```

```
##      sd, var
RLMstep <- function(y, x){
  set.seed(3)
  y <- as.matrix(y) #convert data.frame to a matrix
  x <- as.matrix(x)
  p <- ncol(x) # maximum number of predictors
  indicies <- numeric(p) #store indicies
  model.in <- NULL # model in
  stepwise <- wle::wle.stepwise(y ~ x, num.sol = 3,
                                min.weight = 0.5, type = "Stepwise", method = "WLS")
  #stepwise function taken from the correspodng package
  wstep <- stepwise$wstep #This extract the class item from last iteration

  for (i in 1:p) {
    if (wstep[[i + 1]] == 1) {
      indicies[i] <- i
    }
  }
  model.in <- x[, c(indicies), drop = FALSE]
  print(colnames(model.in))
  rlm.step <- MASS::rlm(y ~ model.in) #new robust regression based on selected variables,
  print(rlm.step)
}

#This is an example
data(Boston)
Boston <- as.matrix(Boston)
x <- Boston[, -14, drop = FALSE]
y <- Boston[, 14, drop = FALSE]
a <- RLMstep(y, x)

## [1] "crim"  "zn"    "indus"
## Call:
## rlm(formula = y ~ model.in)
## Converged in 7 iterations
##
## Coefficients:
## (Intercept)  model.incrim  model.inzn model.inindus
## 26.16509677 -0.24952611  0.06077048 -0.40645237
##
## Degrees of freedom: 506 total; 502 residual
## Scale estimate: 5.46
```

I fail to put the BCa function inside my RLMstep function, instead I write a separate function to estimate the BCa confidence intervals. Again I use Boston data as an example.

```

# For the BCa confidence interval
library(MASS)
library(boot)

data(Boston)
boot.huber <- function(data, indices, maxit=20){
  data <- data[indices,]
  mod <- rlm(medv ~ ., data = data, maxit = maxit)
  coefficients(mod)}

set.seed(1) #set seed
Boston.boot <- boot(data = Boston, statistic = boot.huber, R = 1000, maxit = 100)
# BCa for the first three variables
for (i in 2:4) {
  boot.ci <- boot.ci(Boston.boot, index = i, type = "bca")
  print(boot.ci)
}

```

```

## BOOTSTRAP CONFIDENCE INTERVAL CALCULATIONS
## Based on 1000 bootstrap replicates
##
## CALL :
## boot.ci(boot.out = Boston.boot, type = "bca", index = i)
##
## Intervals :
## Level      BCa
## 95%      (-0.1544, -0.0135 )
## Calculations and Intervals on Original Scale
## BOOTSTRAP CONFIDENCE INTERVAL CALCULATIONS
## Based on 1000 bootstrap replicates
##
## CALL :
## boot.ci(boot.out = Boston.boot, type = "bca", index = i)
##
## Intervals :
## Level      BCa
## 95%      ( 0.0075,  0.0611 )
## Calculations and Intervals on Original Scale
## BOOTSTRAP CONFIDENCE INTERVAL CALCULATIONS
## Based on 1000 bootstrap replicates
##
## CALL :
## boot.ci(boot.out = Boston.boot, type = "bca", index = i)
##

```

```
## Intervals :
## Level      BCa
## 95%      (-0.0775,  0.0768 )
## Calculations and Intervals on Original Scale
```

2 Question 2

Compare your function versus the stepAIC function (direction= both“) using the HtVol data from the first assignment. The set of possible predictors includes Male; Age; Ht; Wt; BMI; BSA, all quadratic terms and all bivariate interactions. To do so I expanded the X variable matrix. In this particular dataset, Male and CT is a bivariate dummy variable, no need for quadratic transformation. To make a more comparable comparison, I rescaled the data structure, other numeric variables are scaled to Normal(0,1).

Load the data first.

```
mydata <- read.csv("https://raw.githubusercontent.com/haowang666/Computational-Stats/master/data/HtVol.csv")

#delete NA
mydata <- na.omit(mydata)

# rescale data
library(dplyr)
mydata <- mydata %>% mutate_each_(funs(scale(.)) %>% as.vector),
                                vars = c("HtVol", "Age", "Ht", "Wt", "BMI", "BSA"))
summary(mydata)
```

| ## | HtVol | Male | CT | Age |
|-------------|-----------|-------------------|------------------|------------------|
| ## Min. | :-1.76626 | Min. :0.0000 | Min. :0.000 | Min. : -2.1887 |
| ## 1st Qu.: | :-0.81265 | 1st Qu.:0.0000 | 1st Qu.:0.000 | 1st Qu.: -0.7048 |
| ## Median : | 0.01854 | Median :1.0000 | Median :1.000 | Median : 0.2692 |
| ## Mean : | 0.00000 | Mean :0.6034 | Mean :0.569 | Mean : 0.0000 |
| ## 3rd Qu.: | 0.60575 | 3rd Qu.:1.0000 | 3rd Qu.:1.000 | 3rd Qu.: 0.7068 |
| ## Max. : | 3.35764 | Max. :1.0000 | Max. :1.000 | Max. : 3.0772 |
| ## | Ht | Wt | BMI | BSA |
| ## Min. | :-2.8745 | Min. : -1.69692 | Min. : -1.2874 | Min. : -2.1852 |
| ## 1st Qu.: | :-0.7380 | 1st Qu.: -0.91895 | 1st Qu.: -0.8527 | 1st Qu.: -0.9270 |
| ## Median : | 0.3443 | Median : 0.02612 | Median : -0.1888 | Median : 0.1880 |
| ## Mean : | 0.0000 | Mean : 0.00000 | Mean : 0.0000 | Mean : 0.0000 |
| ## 3rd Qu.: | 0.6606 | 3rd Qu.: 0.67969 | 3rd Qu.: 0.6076 | 3rd Qu.: 0.7199 |
| ## Max. : | 1.4172 | Max. : 3.17885 | Max. : 3.3252 | Max. : 2.3158 |

```
y <- mydata$HtVol
x <- select(mydata, 2:8)
```

```

#expand the matrix
x <- model.matrix(~(Age+Ht+Wt+BMI+BSA)^2 -1,x)
newdata <- cbind(y,x)
newdata <- as.data.frame(newdata)
newdata$Age2 <- (newdata$Age)^2
newdata$Ht2 <- (newdata$Ht)^2
newdata$Wt2 <- (newdata$Wt)^2
newdata$BMI2 <- (newdata$BMI)^2
newdata$BSA2 <- (newdata$BSA)^2

#check variables
names(newdata)

## [1] "y"      "Age"    "Ht"     "Wt"     "BMI"    "BSA"    "Age:Ht"
## [8] "Age:Wt" "Age:BMI" "Age:BSA" "Ht:Wt"  "Ht:BMI" "Ht:BSA" "Wt:BMI"
## [15] "Wt:BSA" "BMI:BSA" "Age2"    "Ht2"    "Wt2"    "BMI2"   "BSA2"

#reorganize into matrix
y <- newdata$y
y <- as.matrix(y)
x <- newdata[, -1, drop = FALSE]
x <- as.matrix(x)

# get results from RLMstep
time1 <- system.time(RLMstep(y, x))

## [1] "Ht:Wt"  "Ht:BMI"
## Call:
## rlm(formula = y ~ model.in)
## Converged in 7 iterations
##
## Coefficients:
## (Intercept)  model.inHt:Wt  model.inHt:BMI
## 0.19595003   -0.43449973    0.02639283
##
## Degrees of freedom: 58 total; 55 residual
## Scale estimate: 0.769

time1

##      user  system elapsed
##    0.21    0.07    0.27

```

We can see that RLMstep picks two interaction terms. To obtain the BCa CI, I use the following codes. The BCa confidence interval for Ht:Wt is (-0.7597, 0.2413), and BCa CI for Ht:BMI is (-0.8515, 0.5947).

```

# BCa interval
RLM.data <- newdata[, c("y", "Ht:Wt", "Ht:BMI")]
colnames(RLM.data) <- c("y", "HtWt", "HtBMI")

boot.huber <- function(data, indices, maxit=20){
  data <- data[indices,]
  mod <- rlm(y ~ ., data = data, maxit = maxit)
  coefficients(mod)}

set.seed(1) #set seed
RLM.boot <- boot(data = RLM.data, statistic = boot.huber, R = 1000, maxit = 100)
# BCa for the first three variables
for (i in 2:3) {
  boot.ci <- boot.ci(RLM.boot, index = i, type = "bca")
  print(boot.ci)
}

```

```

## BOOTSTRAP CONFIDENCE INTERVAL CALCULATIONS
## Based on 1000 bootstrap replicates
##
## CALL :
## boot.ci(boot.out = RLM.boot, type = "bca", index = i)
##
## Intervals :
## Level      BCa
## 95%      (-0.7597,  0.2413 )
## Calculations and Intervals on Original Scale
## Some BCa intervals may be unstable
## BOOTSTRAP CONFIDENCE INTERVAL CALCULATIONS
## Based on 1000 bootstrap replicates
##
## CALL :
## boot.ci(boot.out = RLM.boot, type = "bca", index = i)
##
## Intervals :
## Level      BCa
## 95%      (-0.8515,  0.5947 )
## Calculations and Intervals on Original Scale

```

The codes for stepAIC is in the following block

```

step$anova # display results

```

```

## Stepwise Model Path
## Analysis of Deviance Table
##

```

```
## Initial Model:
## y ~ Age + Ht + Wt + BMI + BSA + `Age:Ht` + `Age:Wt` + `Age:BMI` +
##      `Age:BSA` + `Ht:Wt` + `Ht:BMI` + `Ht:BSA` + `Wt:BMI` + `Wt:BSA` +
##      `BMI:BSA` + Age2 + Ht2 + Wt2 + BMI2 + BSA2
##
## Final Model:
## y ~ BMI + BSA + `Age:Ht` + `Age:BSA` + `Ht:Wt` + `Ht:BMI` + `Ht:BSA` +
##      `Wt:BMI` + `Wt:BSA` + `BMI:BSA` + Ht2 + Wt2 + BSA2
##
##
##           Step Df      Deviance Resid. Df Resid. Dev      AIC
## 1
## 2 - `Age:Wt`  1 2.153155e-06      38  2.830789 -135.1535
## 3      - Wt  1 7.944901e-05      39  2.830869 -137.1518
## 4      - Age2 1 2.121883e-03      40  2.832991 -139.1084
## 5      - Age  1 3.182429e-03      41  2.836173 -141.0433
## 6      - Ht   1 3.244171e-03      42  2.839417 -142.9770
## 7 - `Age:BMI` 1 3.306997e-03      43  2.842724 -144.9095
## 8      - BMI2 1 6.769789e-02      44  2.910422 -145.5444
```

3 Question 3

Discuss your findings; create a few figures that convey useful info wrt your results.

The major difference between my function and the stepAIC function is the number of predictors. RLMstep only includes two interaction terms, but stepAIC picks the following:

- BMI + BSA + Age:Ht + Age:BSA + Ht:Wt + Ht:BMI + Ht:BSA + Wt:BMI + Wt:BSA + BMI:BSA + Ht2 + Wt2 + BSA2

There could be multiple reasons: stepAIC cannot handle robust regression, thus the two models evaluated are actually different. function `lm` is used in stepAIC, but function `rlm` is used in stepRLM, thus should use weighted AIC


```
wle.aic(formula = y ~ ., data = RLM.data)
```

```
##
## Call:
## wle.aic(formula = y ~ ., data = RLM.data)
##
##
## Weighted Akaike Information Criterion (WAIC):
##      (Intercept) HtWt HtBMI  waic
## [1,]           1    1    0 119.6
## [2,]           0    1    0 121.0
```

```
## [3,]          1      1      1 121.1
## [4,]          0      0      1 121.7
## [5,]          0      1      1 121.8
## [6,]          1      0      1 123.0
## [7,]          1      0      0 139.4
##
## Printed the first 7 best models
```

First I show the residual-fitted plots of the two. Other than creating the resid-fitted plot of the lm function, I also get the resid-fitted plot of the rlm function.

```
rlm <- rlm(y ~ ., data = RLM.data)
residual <- rlm$residuals
fitted <- rlm$fitted.values
plot(fitted, residual)
```




```
#-----
stepAIC.data <- newdata[, c("y", "BMI", "BSA", "Age:Ht", "Age:BSA", "Ht:Wt", "Ht:BMI",
colnames(stepAIC.data) <- c("y", "BMI", "BSA", "AgeHt", "AgeBSA", "HtWt", "HtBMI", "HtBSA",
names(stepAIC.data)
```


```
## [1] "y"      "BMI"    "BSA"    "AgeHt"  "AgeBSA" "HtWt"   "HtBMI"
## [8] "HtBSA"  "WtBMI"  "WtBSA"  "BMIBSA" "Ht2"    "Wt2"    "BSA2"
```



```
lm <- lm(y ~., data = stepAIC.data)
residual <- lm$residuals
fitted <- lm$fitted.values
plot(fitted, residual)
```



```
rlm2 <- rlm(y ~., data = stepAIC.data)
residual <- rlm2$residuals
fitted <- rlm2$fitted.values
plot(fitted, residual)
```



Then I perform another bootstrap for both models (rlm for RLMstep, lm for stepAIC)

```
library(caret)
RMSE_rlm <- function(data, i){
  #index data for resampling
  train_data <- data[i,]
  test_data <- data
  model <- MASS::rlm(y ~., data = train_data)
  predict <- predict(model, newdata = test_data)
  #return rmse
  RMSE <- RMSE(predict, test_data$y)
  return(RMSE)
}

RMSE_step <- function(data, i){
  #index data for resampling
  train_data <- data[i,]
  test_data <- data
  model <- lm(y~., data = train_data)
  predict <- predict(model, newdata = test_data)
  #return rmse
  RMSE <- RMSE(predict, test_data$y)
```

```

return(RMSE)
}

#RLM.data, stepAIC.data

# Perform Bootstrap
Repeats <- 100
set.seed(1)
res <- boot(RLM.data, statistic = RMSE_rlm, R = Repeats)

## Warning in rlm.default(x, y, weights, method = method, wt.method =
## wt.method, : 'rlm' failed to converge in 20 steps

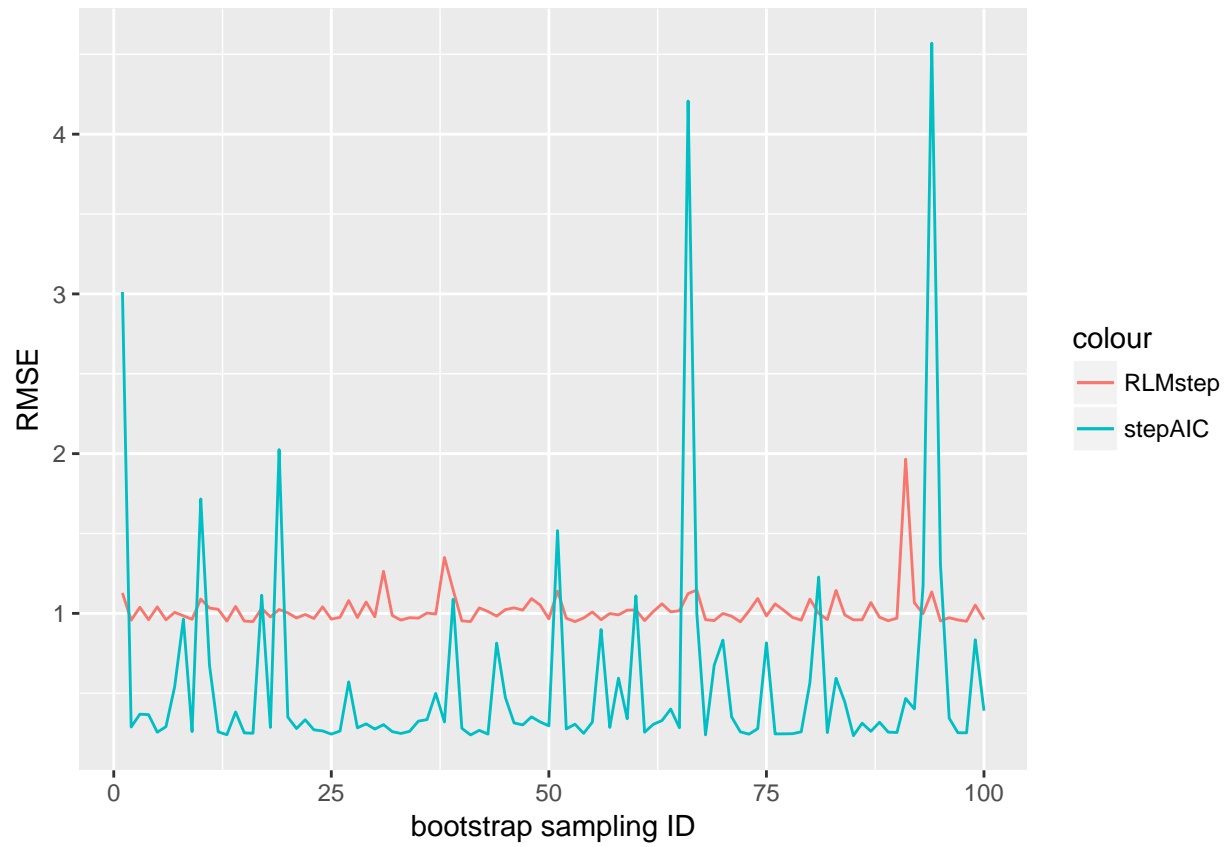
RMSE_rlm <- res$t

set.seed(1)
res <- boot(stepAIC.data, statistic = RMSE_step, R = Repeats)
RMSE_step <- res$t

x <- seq(1:100)
RMSE_res <- cbind(x, RMSE_rlm, RMSE_step)
RMSE_res <- as.data.frame(RMSE_res)
names(RMSE_res) <- c("ID", "RLMstep", "stepAIC")

ggplot(data = RMSE_res, aes(x = ID)) +
  geom_line(aes(y = RLMstep, color = "RLMstep")) +
  geom_line(aes(y = stepAIC, color = "stepAIC")) +
  xlab("bootstrap sampling ID") +
  ylab("RMSE")

```



From the figure, it looks the RLMstep results are more stable.